

Case Study 3

Ronnie Kupfer

2024-07-21

```
library(plsdepot)
library(ggplot2)
library(visreg)
library(caret)
```

```
## Loading required package: lattice
```

```
# Working directory
wd = "C:/Users/ronku/OneDrive/Desktop/WPI Master's Program/DS 501/Case Study 3"
setwd(wd)
library(readxl)
# Create data frame
concrete = read_excel("Concrete_Data.xls")
# Simplify column names without modifying source data file
colnames(concrete)[c(1,2,3,4,5,6,7,8,9)] = c("Cement", "Slag", "Ash", "Water", "Superplasticizer", "C_Aggregate", "F_Aggregate", "Age", "Compressive_Strength")
# View(concrete)
attach(concrete)
# detach(concrete)
# str(concrete)
head(concrete)
```

```
## # A tibble: 6 x 9
##   Cement Slag Ash Water Superplasticizer C_Aggregate F_Aggregate Age
##   <dbl> <dbl> <dbl> <dbl>          <dbl>          <dbl>          <dbl> <dbl>
## 1   540     0     0   162             2.5           1040           676    28
## 2   540     0     0   162             2.5           1055           676    28
## 3   332.  142.     0   228             0             932           594   270
## 4   332.  142.     0   228             0             932           594   365
## 5   199.  132.     0   192             0             978.           826.   360
## 6   266   114     0   228             0             932           670    90
## # i 1 more variable: Compressive_Strength <dbl>
```

R Markdown

This is an R Markdown document. Markdown is a simple formatting syntax for authoring HTML, PDF, and MS Word documents. For more details on using R Markdown see <http://rmarkdown.rstudio.com>.

When you click the **Knit** button a document will be generated that includes both content as well as the output of any embedded R code chunks within the document. You can embed an R code chunk like this:

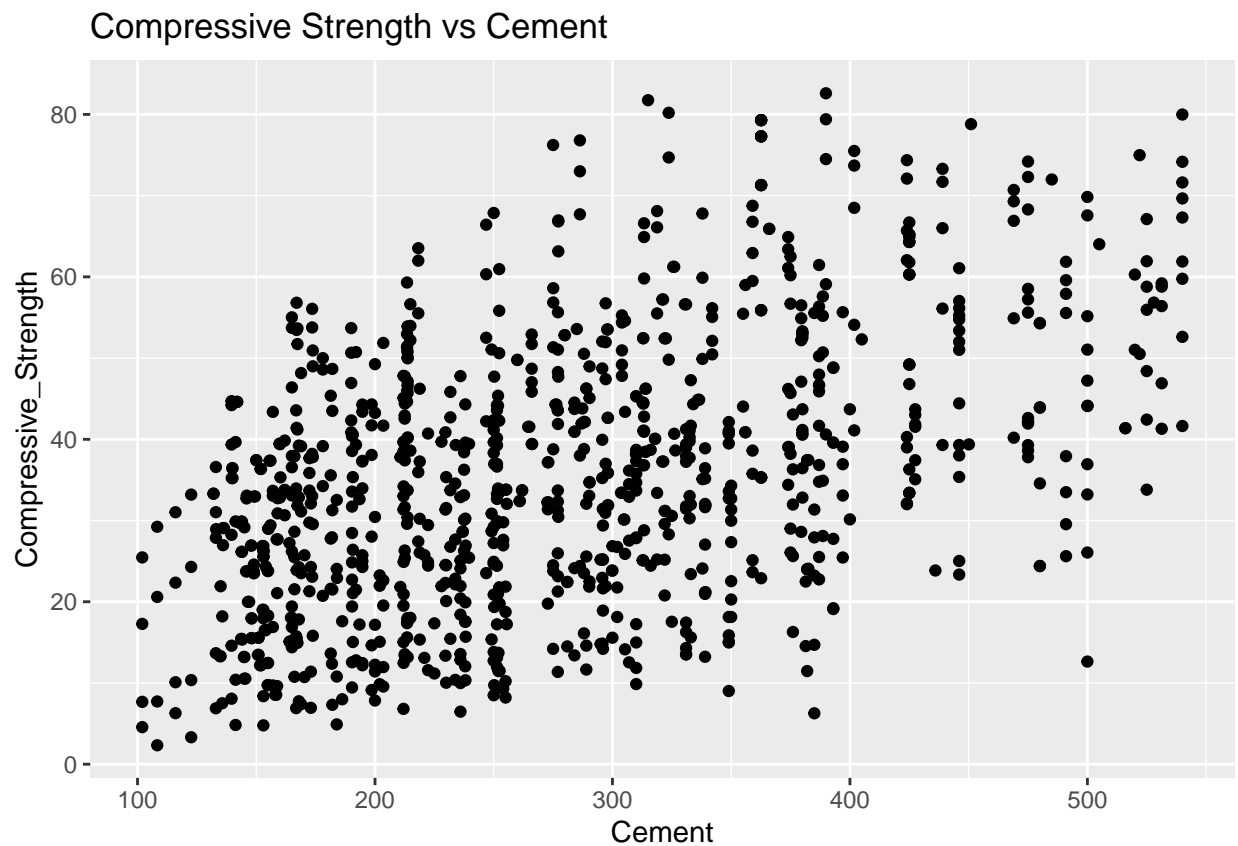
Exploratory Data Analysis -

The author of the posted data stated that there is a strong non-linear relationship between age and compressive strength. The following plots examine the relationship of each input variable to the output variable, compressive strength.

```
strengthVsCement = qplot(Cement, Compressive_Strength, main = "Compressive Strength vs Cement")
```


```
## Warning: 'qplot()' was deprecated in ggplot2 3.4.0.  
## This warning is displayed once every 8 hours.  
## Call 'lifecycle::last_lifecycle_warnings()' to see where this warning was  
## generated.
```

```
strengthVsCement
```

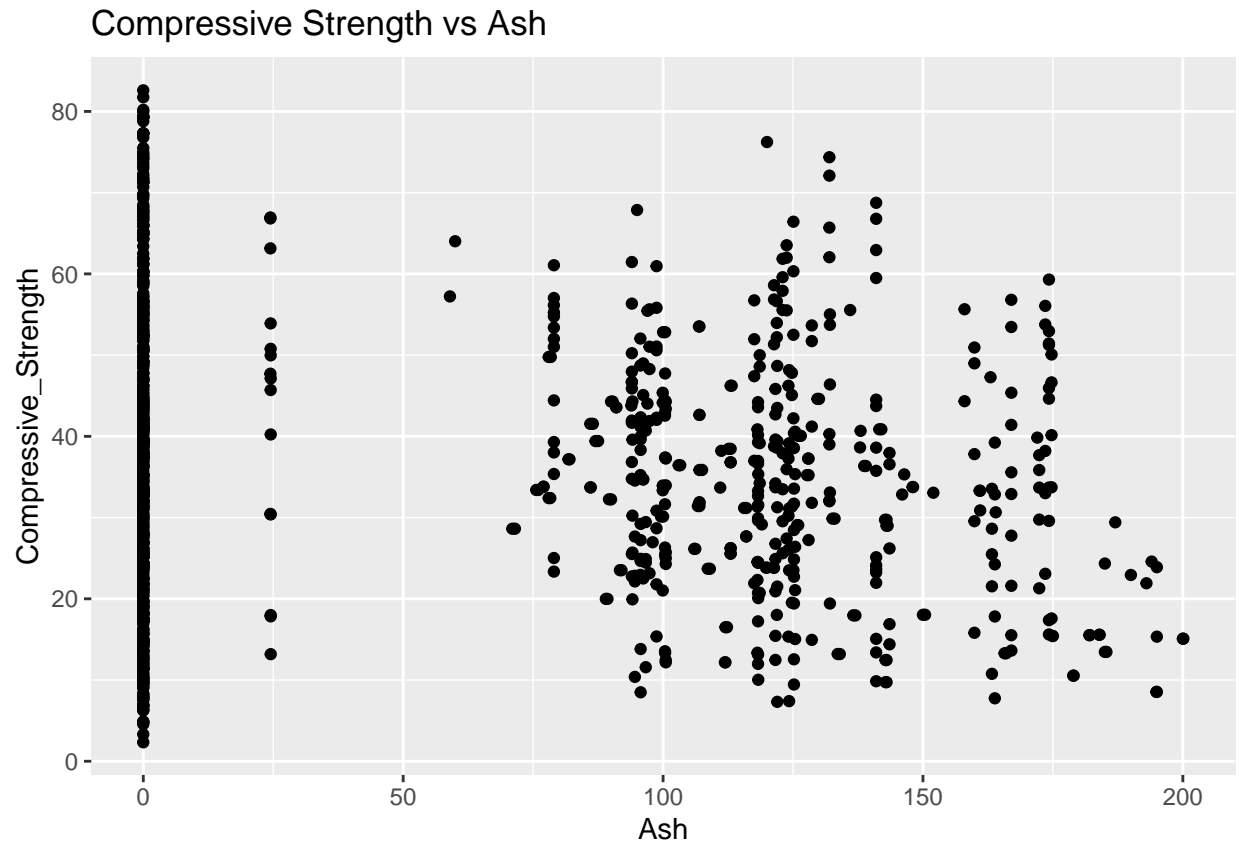


```
strengthVsSlag = qplot(Slag, Compressive_Strength, main = "Compressive Strength vs Slag")  
strengthVsSlag
```


Compressive Strength vs Slag



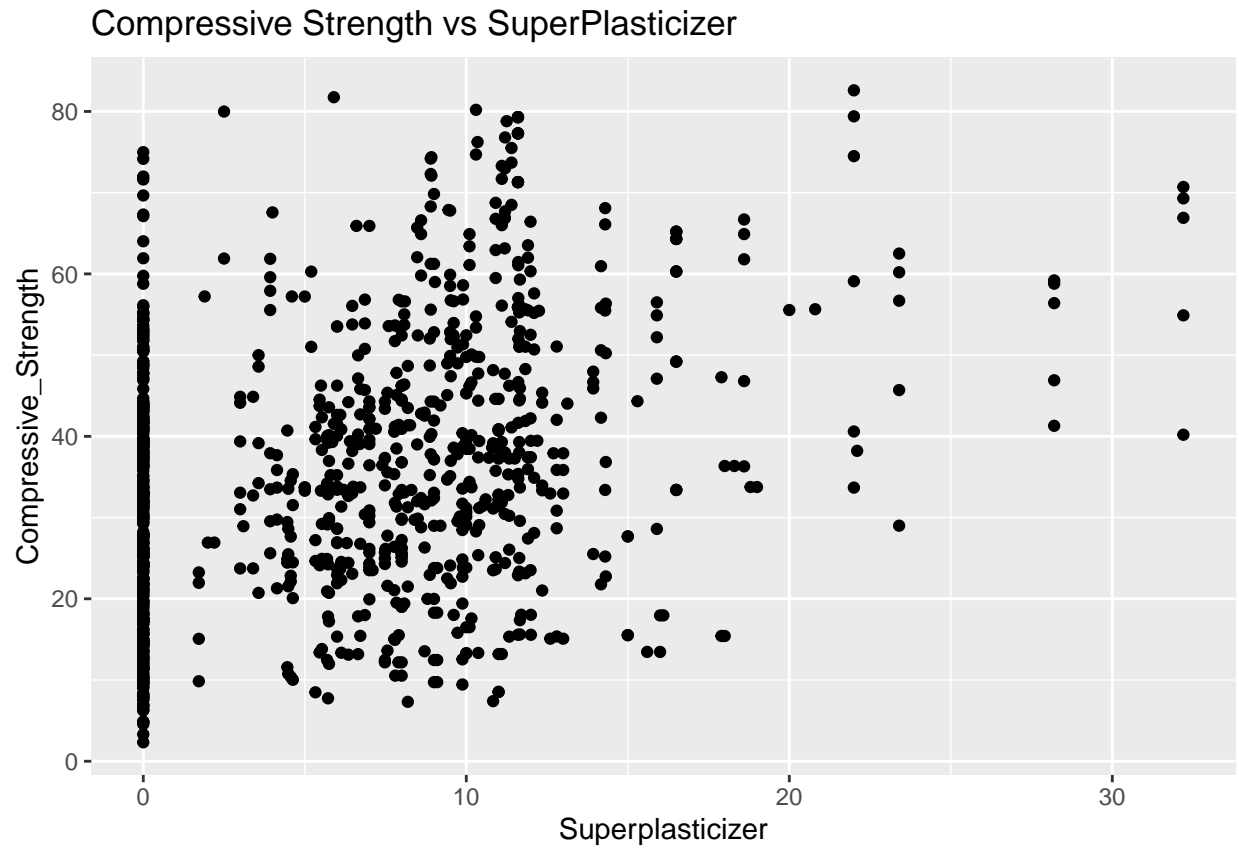
```
strengthVsAsh = qplot(Ash, Compressive_Strength, main = "Compressive Strength vs Ash")
strengthVsAsh
```



```
strengthVsWater = qplot(Water, Compressive_Strength, main = "Compressive Strength vs Water")
strengthVsWater
```




```
strengthVsPlasticizer = qplot(Superplasticizer, Compressive_Strength, main = "Compressive Strength vs S  
strengthVsPlasticizer
```




```
strengthVsC_Aggregate = qplot(C_Aggregate, Compressive_Strength, main = "Compressive Strength vs Course  
strengthVsC_Aggregate
```

Compressive Strength vs Course Aggregate

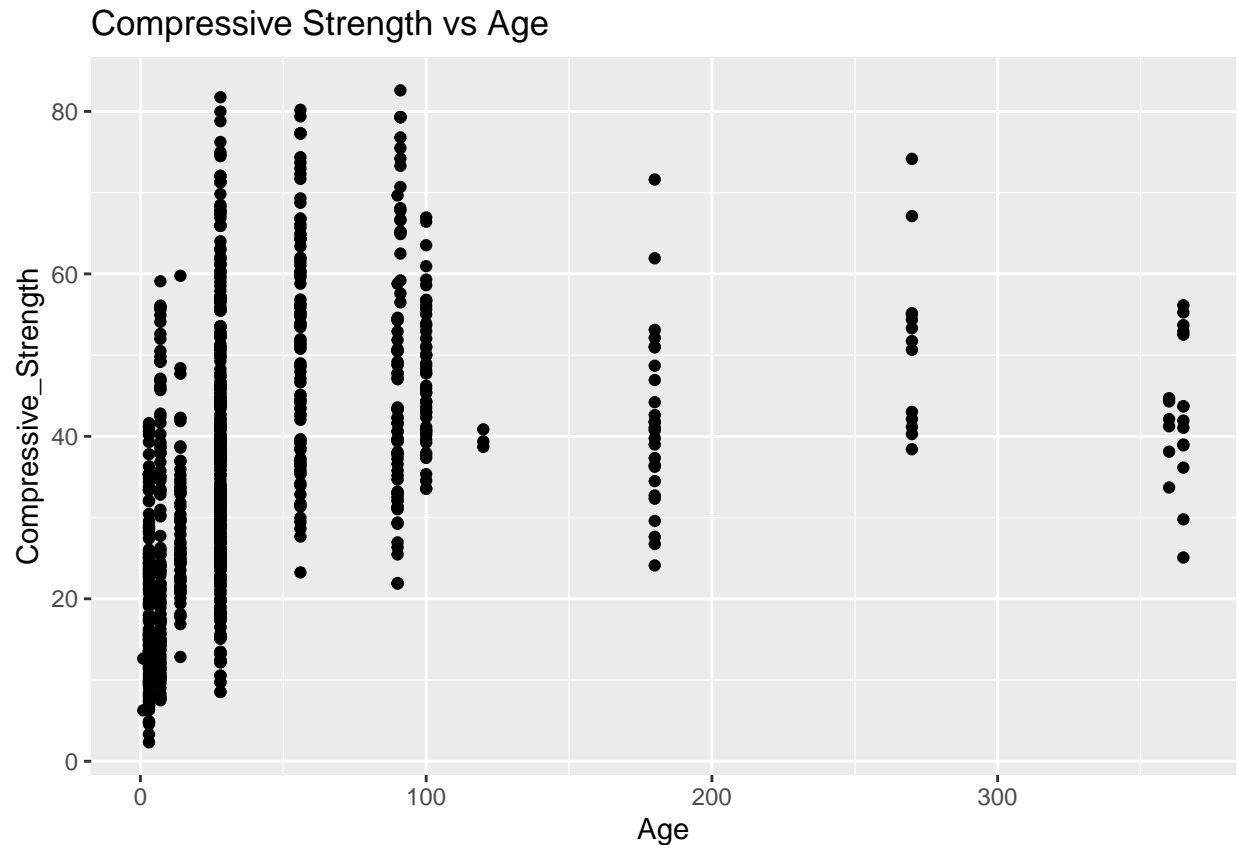


```
strengthVsF_Aggregate = qplot(F_Aggregate, Compressive_Strength, main = "Compressive Strength vs Fine A  
strengthVsF_Aggregate
```

Compressive Strength vs Fine Aggregate



```
strengthVsAge = qplot(Age, Compressive_Strength, main = "Compressive Strength vs Age")
strengthVsAge
```

```
# {r scatter plot Compressive_Strength to Age, echo=T, eval=T, error=TRUE} #plot(concrete)
```

Reviewing scatter plots each of the input variables to the output variable, Compressive Strength, once can draw three conclusions: 1) It appears that as the amount of cement added to the mixture increases the compressive strength of the concrete increases in a fairly linear manner. 2) The data provider's statement that there is a highly nonlinear function of concrete compressive strength to age and ingredients seems to have merit. 3) Most of the input variables seem to have optimal ranges.

```
corStrengthToCement = cor(Compressive_Strength, Cement)
corStrengthToCement
```

```
## [1] 0.4978327
```

```
corStrengthToSlag = cor(Compressive_Strength, Slag)
corStrengthToSlag
```

```
## [1] 0.1348244
```

```
corStrengthToAsh = cor(Compressive_Strength, Ash)
corStrengthToAsh
```

```
## [1] -0.1057533
```

```
corStrengthToWater = cor(Compressive_Strength, Water)
corStrengthToWater
```

```
## [1] -0.2896135
```

```
corStrengthToPlasticizer = cor(Compressive_Strength, Superplasticizer)
corStrengthToPlasticizer
```

```
## [1] 0.3661023
```

```
corStrengthToC_Aggregate = cor(Compressive_Strength, C_Aggregate)
corStrengthToC_Aggregate
```

```
## [1] -0.1649278
```

```
corStrengthToF_Aggregate = cor(Compressive_Strength, F_Aggregate)
corStrengthToF_Aggregate
```

```
## [1] -0.167249
```

```
corStrengthToAge = cor(Compressive_Strength, Age)
corStrengthToAge
```

```
## [1] 0.328877
```

The correlation between the strength of the concrete to the amount of cement is nearly .5. This is the strongest correlation of input variable to the output variable and reflects the trend in the graph. Several of the input variables have a slightly negative correlation. If the relationship is non-linear, the negative correlations may not indicate anything about dependence of the variables.

Split the data as Training and Test sets

```
set.seed(2024)
concreteDF = data.frame(concrete)
splitConcrete = caret::createDataPartition(concreteDF[,1], p = 0.8, list=F, times=1)
#splitConcrete
trainConcrete = concreteDF[splitConcrete,]
head(trainConcrete)
```

```
##   Cement  Slag  Ash  Water  Superplasticizer  C_Aggregate  F_Aggregate  Age
## 1  540.0   0.0   0    162                2.5         1040         676   28
## 2  540.0   0.0   0    162                2.5         1055         676   28
## 3  332.5 142.5   0    228                0.0          932         594  270
## 4  332.5 142.5   0    228                0.0          932         594  365
## 6  266.0 114.0   0    228                0.0          932         670   90
## 7  380.0  95.0   0    228                0.0          932         594  365
##   Compressive_Strength
## 1              79.98611
## 2              61.88737
## 3              40.26954
## 4              41.05278
## 6              47.02985
## 7              43.69830
```

```
testConcrete = concreteDF[-splitConcrete,]
#testConcrete
```

Generating a multiple linear regression model


```
mlr = lm(Compressive_Strength ~ Cement + Slag + Ash + Water + Superplasticizer + C_Aggregate + F_Aggregate + Age, data = trainConcrete)
#fitted(mlr)
#resid(mlr)
#mlr
summary(mlr)
```

```
##
## Call:
## lm(formula = Compressive_Strength ~ Cement + Slag + Ash + Water +
##      Superplasticizer + C_Aggregate + F_Aggregate + Age, data = trainConcrete)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -28.851  -6.312   0.474   6.793  34.164
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   -53.950023   29.053197  -1.857  0.063680 .
## Cement         0.126592    0.009162  13.817 < 2e-16 ***
## Slag           0.109257    0.010969   9.960 < 2e-16 ***
## Ash            0.093200    0.013580   6.863 1.33e-11 ***
## Water        -0.110109    0.044433  -2.478  0.013410 *
## Superplasticizer 0.384435    0.105033   3.660  0.000268 ***
## C_Aggregate    0.029932    0.010271   2.914  0.003663 **
## F_Aggregate    0.031474    0.011763   2.676  0.007606 **
## Age            0.111409    0.005931  18.785 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 10.34 on 817 degrees of freedom
## Multiple R-squared:  0.6227, Adjusted R-squared:  0.619
## F-statistic: 168.6 on 8 and 817 DF, p-value: < 2.2e-16
```

```
#plot(mlr)
```

```
p1 = strengthVsCement + geom_abline(intercept = mlr[1]$coefficients[1], slope = mlr[1]$coefficients[2],
p1
```

Compressive Strength vs Cement




```
predict(mlr, newdata=testConcrete)
```

##	5	9	46	53	57	59	62	66
##	59.88984	19.17717	27.62453	33.57783	68.33219	36.76999	46.13804	47.45548
##	68	75	76	77	80	81	86	95
##	25.97773	55.90629	53.07744	67.72511	62.50311	54.66653	53.10547	49.71150
##	120	125	127	134	148	149	152	153
##	58.37663	50.87755	57.45174	51.10031	53.99699	68.40777	56.81088	57.20757
##	156	159	168	173	175	181	182	185
##	57.20757	59.75504	62.88139	64.47048	60.71018	63.65434	70.65302	20.75366
##	193	194	199	203	206	215	221	228
##	25.83579	30.73776	34.47909	31.17964	24.30879	25.95297	22.38665	38.45155
##	234	240	243	251	255	261	270	273
##	36.49260	21.95076	27.85541	25.41017	27.05262	28.80774	26.60648	32.51113
##	275	291	295	296	299	309	314	317
##	25.31937	27.64829	28.09407	29.31956	38.90070	44.91364	43.69221	37.44771
##	319	322	331	333	337	338	342	345
##	45.46912	34.96442	37.08158	41.76073	37.55365	40.67309	36.48316	32.55770
##	348	354	365	368	369	370	373	374
##	38.46235	43.82052	32.31282	38.21748	43.11945	29.86856	35.77321	40.67519
##	381	386	389	395	401	402	404	411
##	52.93410	50.77951	52.07427	36.56795	46.38364	47.19946	45.52763	27.16347
##	413	423	425	430	437	441	452	456
##	30.69545	32.26849	30.65629	33.35980	29.94869	27.82555	36.60011	38.03895
##	463	467	471	477	481	488	490	491

```
## 37.97010 35.84696 35.97861 50.81969 51.26532 50.95401 49.06068 48.16879
##      494      508      511      512      520      537      538      547
## 48.61443 49.42031 47.08073 52.53975 32.70419 30.65878 39.23035 24.57594
##      551      556      561      568      580      585      586      587
## 28.31920 34.33924 27.76032 23.08091 25.42049 42.45393 38.64528 32.41232
##      589      596      598      599      602      606      610      626
## 19.99638 22.33596 35.91498 21.16242 23.94763 11.72711 31.55653 28.22016
##      630      640      644      648      654      655      656      660
## 14.91512 33.34386 26.31996 33.18096 20.76243 16.23765 26.21408 19.92614
##      670      678      681      682      687      693      707      713
## 33.16740 11.51553 13.85510 17.90894 33.61541 28.68604 21.71106 31.82954
##      720      721      723      727      730      732      735      740
## 33.12141 41.45193 21.34667 23.66531 33.35785 26.49416 36.18670 22.98939
##      741      744      750      754      771      775      780      781
## 18.17822 33.17696 41.85049 51.44683 26.13035 32.63737 22.34434 17.22476
##      782      783      784      792      800      804      812      815
## 19.71360 22.05318 28.96051 45.84966 48.66162 26.98692 23.05860 60.04623
##      826      828      830      848      849      854      859      869
## 33.28707 51.85432 31.64093 42.58063 31.29911 29.32003 39.58656 39.02344
##      883      885      893      902      903      905      913      916
## 22.80964 28.02528 26.28366 30.00342 40.76592 24.56357 36.43688 32.66100
##      918      919      923      924      930      935      938      942
## 24.95739 17.05611 32.02042 33.06891 25.56601 25.65504 34.44796 22.07644
##      947      948      951      954      958      966      967      968
## 32.80925 27.99132 34.51552 16.73157 30.52072 40.71782 17.56750 24.58391
##      976      979      982      986      987      993      994     1005
## 36.46145 32.76523 17.04413 32.28224 31.63149 27.88523 31.30673 42.72968
##     1014     1020     1025     1026
## 38.85186 19.46307 34.99607 39.12007
```

```
predCompressive_Strength = data.frame(predict(mlr, newdata=testConcrete))
names(predCompressive_Strength)[1] = 'Predicted'
predCompressive_Strength$Reference = testConcrete[,c('Compressive_Strength')]
plotpredmlr = qplot(Reference, Predicted, data=predCompressive_Strength) + geom_point(colour = "blue", size = 10)
plotpredmlr
```



Model evaluation - $RMSEP$ and R^2 * Calculating predicted residual sum of squares (PRESS)

$$PRESS = \sum_{i=1}^n (y_i^{ref} - y_i^{pred})^2$$

```
PRESS = sum((predCompressive_Strength$Reference - predCompressive_Strength$Predicted)^2)
PRESS
```

```
## [1] 23498.91
```

- Root mean squared error of prediction (RMSEP)

$$RMSEP = \sqrt{\frac{1}{n_T} \sum_{i=1}^{n_T} (y_i^{ref} - y_i^{pred})^2}$$

```
RMSEP = sqrt(PRESS / nrow(predCompressive_Strength))
RMSEP
```

```
## [1] 10.7327
```

- Total sum of squares (SST)

$$SST = \sum_{i=1}^n (y_i^{ref} - y_i^{mean})^2$$

```
SST = sum((predCompressive_Strength$Reference - mean(predCompressive_Strength$Reference))^2)
SST
```

```
## [1] 55712.38
```

- Calculating R^2

$$R^2 = 1 - \frac{PRESS}{SST}$$


```
R2 = 1 - (PRESS/SST)
R2
```

```
## [1] 0.5782102
```

Predicted versus Reference

- Confidence Intervals (Narrow)

```
predCompressive_Strength$lower = predict(mlr, newdata=testConcrete, interval = "confidence")[,2]
predCompressive_Strength$upper = predict(mlr, newdata=testConcrete, interval = "confidence")[,3]
#predCompressive_Strength
qplot(Reference, Predicted, data=predCompressive_Strength) + geom_point(colour = "blue", size = 3) +
  geom_errorbar(aes(ymin = lower, ymax = upper))
```




* Prediction Intervals (Tolerance/Wide)

```


predCompressive_Strength$lower = predict(mlr, newdata=testConcrete, interval = "prediction")[,2]
predCompressive_Strength$upper = predict(mlr, newdata=testConcrete, interval = "prediction")[,3]
#predCompressive_Strength
qplot(Reference, Predicted, data=predCompressive_Strength) + geom_point(colour = "blue", size = 3) +
  geom_errorbar(aes(ymin = lower,ymax = upper))


```





Visualizing using visreg package

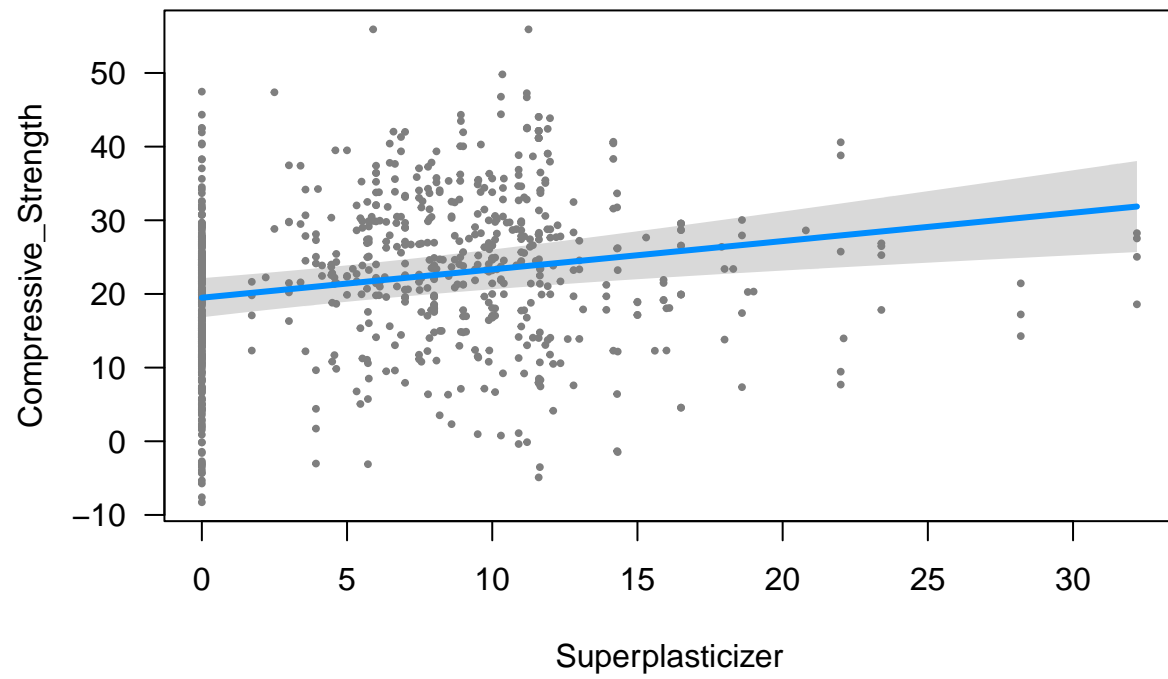
```
visreg::visreg(mlr)
```





















Reviewing the statistics above would indicate that Cement, Slag, Ash, and Age are the most significant input variables in the linear model. The following code creates a new model with only these input variables.

```
mlr2 = lm(Compressive_Strength ~ Cement + Slag + Ash + Age, data=trainConcrete)
#fitted(mlr2)
#resid(mlr2)
#mlr2
summary(mlr2)
```

```
##
## Call:
## lm(formula = Compressive_Strength ~ Cement + Slag + Ash + Age,
##     data = trainConcrete)
##
## Residuals:
```

	Min	1Q	Median	3Q	Max
	-34.949	-7.630	-0.007	7.737	43.322

```
##
## Coefficients:
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-15.619180	1.930572	-8.09	2.13e-15 ***
Cement	0.122720	0.004744	25.87	< 2e-16 ***
Slag	0.094326	0.005549	17.00	< 2e-16 ***
Ash	0.107360	0.008020	13.39	< 2e-16 ***
Age	0.088191	0.006369	13.85	< 2e-16 ***


```
## ---
```



```
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 11.62 on 821 degrees of freedom
## Multiple R-squared:  0.5212, Adjusted R-squared:  0.5189
## F-statistic: 223.4 on 4 and 821 DF,  p-value: < 2.2e-16
```

```
#plot(mlr2)
```

```
p2 = strengthVsCement + geom_abline(intercept = mlr2[1]$coefficients[1], slope = mlr2[1]$coefficients[2])
p2
```




```
predict(mlr2, newdata=testConcrete)
```

```
##      5      9     46     53     57     59     62     66
## 52.99048 30.24674 41.94132 43.86628 74.86247 44.73077 51.58899 53.74696
##      68     75     76     77     80     81     86     95
## 39.24386 39.51304 54.14315 53.25591 49.84640 46.82809 45.47956 48.17841
##     120     125     127     134     148     149     152     153
## 49.03287 43.69808 49.03287 40.93484 46.16743 54.52053 47.56095 51.64523
##     156     159     168     173     175     181     182     185
## 51.64523 59.88911 61.90397 54.58891 50.64764 62.97581 58.08217 22.31186
##     193     194     199     203     206     215     221     228
## 28.16675 32.04716 27.88285 26.18268 24.69781 21.44321 23.52670 31.48833
##     234     240     243     251     255     261     270     273
## 31.30761 25.51733 30.19146 26.56872 21.50153 22.47991 26.14283 30.81696
##     275     291     295     296     299     309     314     317
```

```
## 28.19094 25.31457 22.69222 23.66232 31.24676 37.65524 39.75731 28.48180
##      319      322      331      333      337      338      342      345
## 34.83156 28.02147 29.33501 33.03904 33.63484 36.10419 35.93668 29.63561
##      348      354      365      368      369      370      373      374
## 34.30975 38.10702 29.17870 33.85283 37.73324 24.82922 29.50335 33.38376
##      381      386      389      395      401      402      404      411
## 51.47630 50.66445 48.69825 36.55167 44.70791 45.14206 41.76314 25.58925
##      413      423      425      430      437      441      452      456
## 30.18087 32.21419 30.26347 30.52258 27.79403 23.64799 34.85500 34.22661
##      463      467      471      477      481      488      490      491
## 34.14379 29.99775 36.74548 50.12366 50.47643 46.32108 44.11630 44.11630
##      494      508      511      512      520      537      538      547
## 44.46906 55.13004 53.27803 57.59939 38.25524 35.07904 45.75566 27.71585
##      551      556      561      568      580      585      586      587
## 33.22702 34.90490 27.83141 22.77160 24.62362 45.24702 40.71543 31.80347
##      589      596      598      599      602      606      610      626
## 19.55437 21.40639 37.48593 26.24739 28.45217 13.60726 29.21709 29.99299
##      630      640      644      648      654      655      656      660
## 12.61010 35.93808 27.95007 33.84585 19.26724 19.24208 30.92007 20.92704
##      670      678      681      682      687      693      707      713
## 38.09918 11.94738 13.79939 19.02245 41.30698 31.66281 20.71535 35.72610
##      720      721      723      727      730      732      735      740
## 36.38792 46.08908 23.04128 25.26563 32.93826 27.47459 35.14722 23.17522
##      741      744      750      754      771      775      780      781
## 21.44593 35.56992 46.97537 53.11884 27.47459 31.87711 23.05250 16.05747
##      782      783      784      792      800      804      812      815
## 21.32321 23.17522 28.64307 43.08442 50.91406 32.87426 24.89330 54.17276
##      826      828      830      848      849      854      859      869
## 35.56992 49.05787 39.06037 44.11570 35.08449 39.40242 40.42044 41.32499
##      883      885      893      902      903      905      913      916
## 27.63432 36.54931 30.01507 34.83605 43.50575 23.13330 41.64940 38.58012
##      918      919      923      924      930      935      938      942
## 21.51966 23.86199 35.74866 44.72209 34.04201 37.94102 37.42996 23.01140
##      947      948      951      954      958      966      967      968
## 27.45760 36.50329 41.72313 27.21346 35.69812 43.40666 21.17821 23.14649
##      976      979      982      986      987      993      994      1005
## 41.63053 38.61670 23.90034 35.76661 39.02379 23.25001 37.30403 44.12467
##      1014      1020      1025      1026
## 45.72832 33.16396 31.71799 41.40629
```

```
predCompressive_Strength2 = data.frame(predict(mlr2, newdata=testConcrete))
names(predCompressive_Strength2)[1] = 'Predicted'
predCompressive_Strength2$Reference = testConcrete[,c('Compressive_Strength')]
plotpredmlr2 = qplot(Reference, Predicted, data=predCompressive_Strength2) + geom_point(colour = "blue")
plotpredmlr2
```



Model evaluation - $RMSEP$ and R^2

- Calculating predicted residual sum of squares (PRESS)

$$PRESS = \sum_{i=1}^n (y_i^{ref} - y_i^{pred})^2$$

```
PRESS2 = sum((predCompressive_Strength2$Reference - predCompressive_Strength2$Predicted)^2)
PRESS2
```

```
## [1] 26386.17
```

- Root mean squared error of prediction (RMSEP)

$$RMSEP = \sqrt{\frac{1}{n_T} \sum_{i=1}^{n_T} (y_i^{ref} - y_i^{pred})^2}$$

```
RMSEP2 = sqrt(PRESS2 / nrow(predCompressive_Strength2))
RMSEP2
```

```
## [1] 11.37295
```

- Total sum of squares (SST)

$$SST = \sum_{i=1}^n (y_i^{ref} - y_i^{mean})^2$$

```
SST2 = sum((predCompressive_Strength2$Reference - mean(predCompressive_Strength2$Reference))^2)
SST2
```

```
## [1] 55712.38
```

- Calculating R^2

$$R^2 = 1 - \frac{PRESS}{SST}$$


```
R22 = 1 - (PRESS2/SST2)
R22
```

```
## [1] 0.5263859
```

Predicted versus Reference

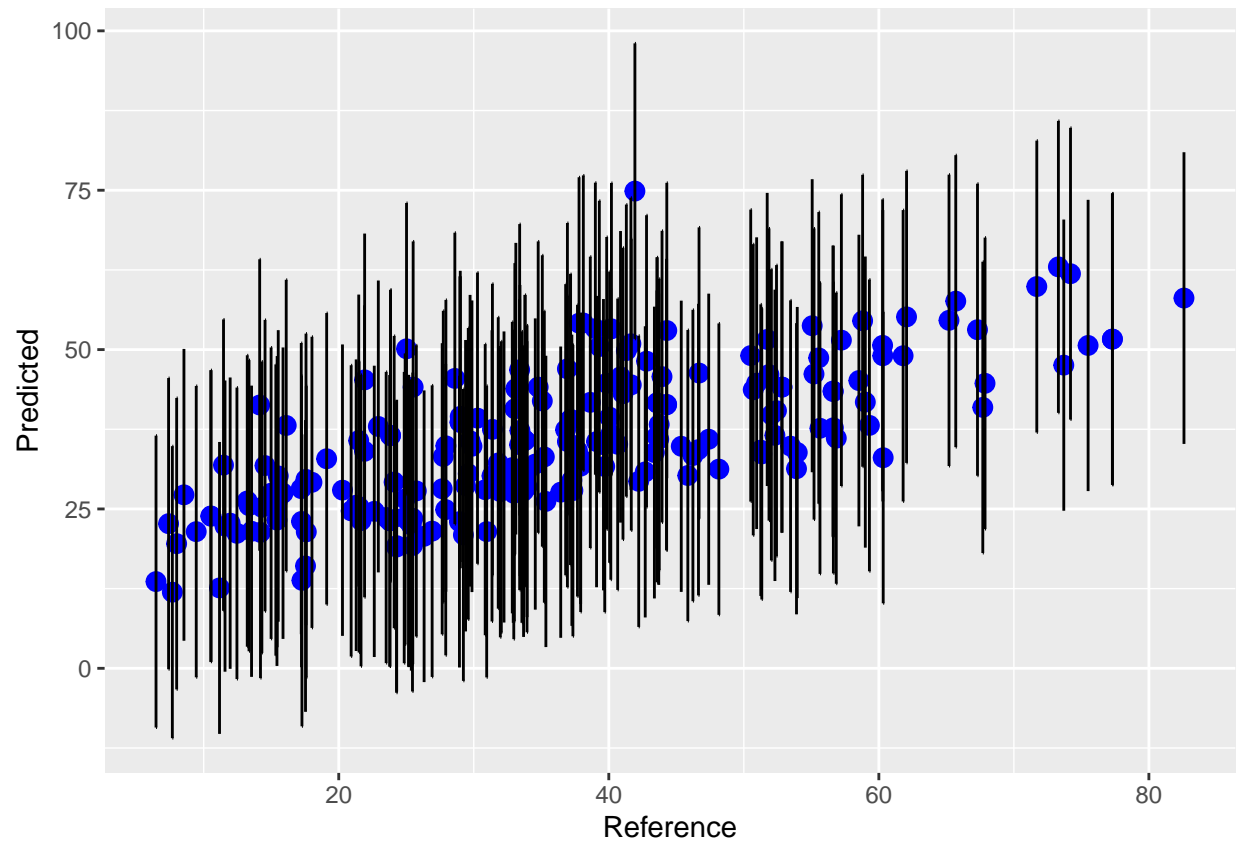
- Confidence Intervals (Narrow)

```
predCompressive_Strength2$lower = predict(mlr2, newdata=testConcrete, interval = "confidence")[,2]
predCompressive_Strength2$upper = predict(mlr2, newdata=testConcrete, interval = "confidence")[,3]
#predCompressive_Strength2
qplot(Reference, Predicted, data=predCompressive_Strength2) + geom_point(colour = "blue", size = 3) +
  geom_errorbar(aes(ymin = lower, ymax = upper))
```




- Prediction Intervals (Tolerance/Wide)


```
predCompressive_Strength2$lower = predict(mlr2, newdata=testConcrete, interval = "prediction")[,2]
predCompressive_Strength2$upper = predict(mlr2, newdata=testConcrete, interval = "prediction")[,3]
#predCompressive_Strength2
qplot(Reference, Predicted, data=predCompressive_Strength2) + geom_point(colour = "blue", size = 3) +
  geom_errorbar(aes(ymin = lower,ymax = upper))
```





Visualizing using `visreg` package

```
visreg::visreg(mlr2)
```









Generating a polynomial multiple linear regression model

```
polymlr = lm(Compressive_Strength ~ poly(Cement + Slag + Ash + Water + Superplasticizer + C_Aggregate +
#fitted(polymlr)
#resid(polymlr)
#polymlr
summary(polymlr)
```

```
##
## Call:
## lm(formula = Compressive_Strength ~ poly(Cement + Slag + Ash +
##      Water + Superplasticizer + C_Aggregate + F_Aggregate + Age,
##      5), data = trainConcrete)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -31.823  -9.835  -0.885   8.943  45.690
##
## Coefficients:
##
## (Intercept)                                     Estimate
## poly(Cement + Slag + Ash + Water + Superplasticizer + C_Aggregate + F_Aggregate + Age, 5)1 290.2800
## poly(Cement + Slag + Ash + Water + Superplasticizer + C_Aggregate + F_Aggregate + Age, 5)2 -57.5966
```


```

## poly(Cement + Slag + Ash + Water + Superplasticizer + C_Aggregate + F_Aggregate + Age, 5)3 -96.6034
## poly(Cement + Slag + Ash + Water + Superplasticizer + C_Aggregate + F_Aggregate + Age, 5)4 43.5343
## poly(Cement + Slag + Ash + Water + Superplasticizer + C_Aggregate + F_Aggregate + Age, 5)5 46.4451
## Std. Error: 0.438
## (Intercept)
## poly(Cement + Slag + Ash + Water + Superplasticizer + C_Aggregate + F_Aggregate + Age, 5)1 12.613
## poly(Cement + Slag + Ash + Water + Superplasticizer + C_Aggregate + F_Aggregate + Age, 5)2 12.613
## poly(Cement + Slag + Ash + Water + Superplasticizer + C_Aggregate + F_Aggregate + Age, 5)3 12.613
## poly(Cement + Slag + Ash + Water + Superplasticizer + C_Aggregate + F_Aggregate + Age, 5)4 12.613
## poly(Cement + Slag + Ash + Water + Superplasticizer + C_Aggregate + F_Aggregate + Age, 5)5 12.613
## t value
## (Intercept) 81.408
## poly(Cement + Slag + Ash + Water + Superplasticizer + C_Aggregate + F_Aggregate + Age, 5)1 23.013
## poly(Cement + Slag + Ash + Water + Superplasticizer + C_Aggregate + F_Aggregate + Age, 5)2 -4.566
## poly(Cement + Slag + Ash + Water + Superplasticizer + C_Aggregate + F_Aggregate + Age, 5)3 -7.659
## poly(Cement + Slag + Ash + Water + Superplasticizer + C_Aggregate + F_Aggregate + Age, 5)4 3.451
## poly(Cement + Slag + Ash + Water + Superplasticizer + C_Aggregate + F_Aggregate + Age, 5)5 3.682
## Pr(>|t|)
## (Intercept) < 2e-16
## poly(Cement + Slag + Ash + Water + Superplasticizer + C_Aggregate + F_Aggregate + Age, 5)1 < 2e-16
## poly(Cement + Slag + Ash + Water + Superplasticizer + C_Aggregate + F_Aggregate + Age, 5)2 5.73e-06
## poly(Cement + Slag + Ash + Water + Superplasticizer + C_Aggregate + F_Aggregate + Age, 5)3 5.29e-14
## poly(Cement + Slag + Ash + Water + Superplasticizer + C_Aggregate + F_Aggregate + Age, 5)4 0.000586
## poly(Cement + Slag + Ash + Water + Superplasticizer + C_Aggregate + F_Aggregate + Age, 5)5 0.000246
## ***
## poly(Cement + Slag + Ash + Water + Superplasticizer + C_Aggregate + F_Aggregate + Age, 5)1 ***
## poly(Cement + Slag + Ash + Water + Superplasticizer + C_Aggregate + F_Aggregate + Age, 5)2 ***
## poly(Cement + Slag + Ash + Water + Superplasticizer + C_Aggregate + F_Aggregate + Age, 5)3 ***
## poly(Cement + Slag + Ash + Water + Superplasticizer + C_Aggregate + F_Aggregate + Age, 5)4 ***
## poly(Cement + Slag + Ash + Water + Superplasticizer + C_Aggregate + F_Aggregate + Age, 5)5 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 12.61 on 820 degrees of freedom
## Multiple R-squared:  0.4363, Adjusted R-squared:  0.4328
## F-statistic: 126.9 on 5 and 820 DF,  p-value: < 2.2e-16


#plot(polymlr)
polyp1 = strengthVsCement + geom_abline(intercept = polymlr[1]$coefficients[1], slope = polymlr[1]$coef
polyp1

```

Compressive Strength vs Cement



```
ggplot(trainConcrete, aes(Water, Compressive_Strength) ) + geom_point() +  
stat_smooth(method = lm, formula = y ~ poly(x, 2, raw = TRUE))
```




```
predict(polymlr, newdata=testConcrete)
```

```
##      5      9     46     53     57     59     62     66
## 41.71701 23.49457 23.43686 24.77183 56.65083 34.04772 52.81526 52.81526
##      68     75     76     77     80     81     86     95
## 23.43686 50.26645 40.36885 47.23459 46.87155 45.44518 45.29260 41.44354
##    120    125    127    134    148    149    152    153
## 50.75817 49.26005 50.66048 52.70797 54.42773 56.41653 57.28469 53.61689
##    156    159    168    173    175    181    182    185
## 53.61689 49.34161 57.07640 59.09037 59.56060 55.56027 59.39423 28.02813
##    193    194    199    203    206    215    221    228
## 33.29687 42.48935 52.92193 43.80159 30.58789 32.35946 29.40020 35.25365
##    234    240    243    251    255    261    270    273
## 52.09433 26.27931 34.22738 30.59766 35.08800 36.91494 30.79942 41.35558
##    275    291    295    296    299    309    314    317
## 28.00084 32.49109 33.00150 35.11788 53.07374 55.48451 53.86999 45.55404
##    319    322    331    333    337    338    342    345
## 57.89908 41.05456 42.23423 51.12438 41.35558 47.46837 36.70709 33.34053
##    348    354    365    368    369    370    373    374
## 44.52247 53.13165 28.09738 37.42578 47.00205 33.93432 45.20749 53.92176
##    381    386    389    395    401    402    404    411
## 45.74969 43.18468 35.62769 40.93824 32.45134 32.28051 34.05350 33.09417
##    413    423    425    430    437    441    452    456
## 29.21381 31.49393 32.56175 33.68609 38.11366 35.60950 39.13878 42.50325
##    463    467    471    477    481    488    490    491
```

```
## 53.16141 50.93511 50.21287 36.99207 37.83330 52.92370 48.52471 48.04022
##      494      508      511      512      520      537      538      547
## 48.87237 29.62293 26.86068 34.52926 23.84496 26.40673 27.42898 25.30485
##      551      556      561      568      580      585      586      587
## 24.72395 35.04621 32.37025 30.07417 33.74127 34.61157 36.01362 49.56500
##      589      596      598      599      602      606      610      626
## 27.71154 30.76644 35.04621 23.73055 24.44262 23.71122 49.36198 30.48887
##      630      640      644      648      654      655      656      660
## 28.81885 33.10160 29.40020 34.35750 29.85589 23.86246 24.63752 27.79965
##      670      678      681      682      687      693      707      713
## 25.15659 23.68998 24.07236 24.42308 24.38021 30.52126 30.18464 25.16381
##      720      721      723      727      730      732      735      740
## 31.96066 35.16574 27.74918 30.81594 48.74646 37.90964 55.42877 30.32795
##      741      744      750      754      771      775      780      781
## 23.69068 33.10160 47.27716 52.63611 35.42612 50.95254 27.14558 28.67828
##      782      783      784      792      800      804      812      815
## 27.03097 29.85589 42.03806 58.68487 47.70121 23.71122 27.62441 42.14214
##      826      828      830      848      849      854      859      869
## 32.91656 50.36544 23.76810 27.38094 24.34788 23.71122 28.00474 26.18811
##      883      885      893      902      903      905      913      916
## 23.67697 23.67023 23.81687 23.81687 27.03097 24.44262 25.03003 23.78306
##      918      919      923      924      930      935      938      942
## 27.87596 23.69861 24.96252 23.69861 23.67427 23.69381 25.67672 24.53421
##      947      948      951      954      958      966      967      968
## 26.63535 23.67014 23.85192 22.75941 23.80100 26.96318 23.67250 24.46746
##      976      979      982      986      987      993      994     1005
## 24.99602 23.77393 23.69772 24.99602 23.75570 24.30325 24.26883 27.50166
##     1014     1020     1025     1026
## 24.28590 22.98643 26.01668 26.24674
```

```
predCompressive_Strength3 = data.frame(predict(polymlr, newdata=testConcrete))
names(predCompressive_Strength3)[1] = 'Predicted'
predCompressive_Strength3$Reference = testConcrete[,c('Compressive_Strength')]
plotpredpolymlr = qplot(Reference, Predicted, data=predCompressive_Strength3) + geom_point(colour = "black")
plotpredpolymlr
```



Model evaluation - $RMSEP$ and R^2

- Calculating predicted residual sum of squares (PRESS)

$$PRESS = \sum_{i=1}^n (y_i^{ref} - y_i^{pred})^2$$

```
PRESS3 = sum((predCompressive_Strength3$Reference - predCompressive_Strength3$Predicted)^2)
PRESS3
```

```
## [1] 38521.65
```

- Root mean squared error of prediction (RMSEP)

$$RMSEP = \sqrt{\frac{1}{n_T} \sum_{i=1}^{n_T} (y_i^{ref} - y_i^{pred})^2}$$

```
RMSEP3 = sqrt(PRESS3 / nrow(predCompressive_Strength3))
RMSEP3
```

```
## [1] 13.7416
```

- Total sum of squares (SST)

$$SST = \sum_{i=1}^n (y_i^{ref} - y_i^{mean})^2$$

```
SST3 = sum((predCompressive_Strength3$Reference - mean(predCompressive_Strength3$Reference))^2)
SST3
```

```
## [1] 55712.38
```

- Calculating R^2

$$R^2 = 1 - \frac{PRESS}{SST}$$


```
R23 = 1 - (PRESS3/SST3)
R23
```

```
## [1] 0.3085622
```

Predicted versus Reference


- Confidence Intervals (Narrow)

```
predCompressive_Strength3$lower = predict(polymlr, newdata=testConcrete, interval = "confidence")[,2]
predCompressive_Strength3$upper = predict(polymlr, newdata=testConcrete, interval = "confidence")[,3]
#predCompressive_Strength3
qplot(Reference, Predicted, data=predCompressive_Strength3) + geom_point(colour = "blue", size = 3) +
  geom_errorbar(aes(ymin = lower, ymax = upper))
```

* Prediction Intervals (Tolerance/Wide)

```
predCompressive_Strength3$lower = predict(polymlr, newdata=testConcrete, interval = "prediction")[,2]
predCompressive_Strength3$upper = predict(polymlr, newdata=testConcrete, interval = "prediction")[,3]
#predCompressive_Strength3
qplot(Reference, Predicted, data=predCompressive_Strength3) + geom_point(colour = "blue", size = 3) +
  geom_errorbar(aes(ymin = lower,ymax = upper))
```



Visualizing using visreg package

```
visreg::visreg(polymlr)
```

